




A Brief Introduction to **Machine Learning** (机器学习简介)


Road Map

- The Concept of Machine Learning
- Learners (学习器)
- Learning Paradigms (学习范式)
- Resources
- Summary




Road Map

- The Concept of Machine Learning
- Learners (学习器)
- Learning Paradigms (学习范式)
- Resources
- Summary



Machine Learning?

- No fixed definition
 - “利用**经验**改善系统自身的性能” [Mitchell, Book97]
- “Experience” is usually expressed as **data** in computers
- Main tasks: to learn knowledge from data and make **predictions**
- DMP: D(Training Data)、M(Model)、P(Prediction)




Play Tennis

Day	Outlook	Temperature	Humidity	Wind	Play Tennis?
1	Sunny	Hot	High	Light	No
2	Sunny	Hot	High	Strong	No
3	Overcast	Hot	High	Light	Yes
4	Rain	Mild	High	Light	Yes
5	Rain	Cool	Normal	Light	Yes
6	Rain	Cool	Normal	Strong	No
7	Overcast	Cool	Normal	Strong	Yes
8	Sunny	Mild	High	Light	No
9	Sunny	Cool	Normal	Light	Yes
10	Rain	Mild	Normal	Light	Yes
11	Sunny	Mild	Normal	Strong	Yes
12	Overcast	Mild	High	Strong	Yes
13	Overcast	Hot	Normal	Light	Yes
14	Rain	Mild	High	Strong	No
15	Sunny	Hot	Normal	Strong	Yes/No?



What to Learn?



given a new instance \mathbf{x} :

--	--	--


 ...

--

predict the label y : ■ or ■ ?

Learn a **decision function**

$$y = f(\mathbf{x})$$




Supervised vs. Unsupervised

■ Supervised learning

- $(\mathbf{X}=\{\mathbf{x}_1, \dots, \mathbf{x}_n\}, \mathbf{Y}=\{y_1, \dots, y_n\})$ known, learn $y = f(\mathbf{x})$
- **Classification**: y_i is a **discrete** class label
- Regression: y_i is a **continuous** value

■ Unsupervised learning

- \mathbf{X} known, \mathbf{Y} unknown
- Clustering analysis
- Anomaly detection



Applications—Machine Learning+



Computer
Vision

Natural Language
Processing

..... Information
Security

Bioinformatics
& Biometrics

Information
Retrieval

..... Robotics &
Automation

application driven


Machine Learning

underpinning and supporting




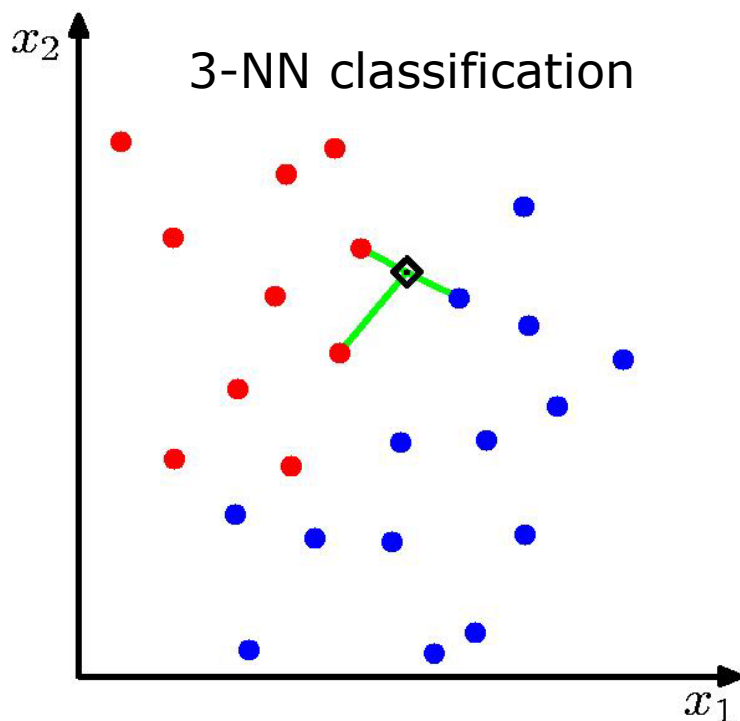
Road Map

- Machine Learning
- Learners
- Learning Paradigms
- Resources
- Conclusion



k -Nearest Neighbors (k -NN)

- For each unknown instance
 - find its k **nearest** neighbors: distance metric
 - pick the class label with the **most votes**



k -NN

特点

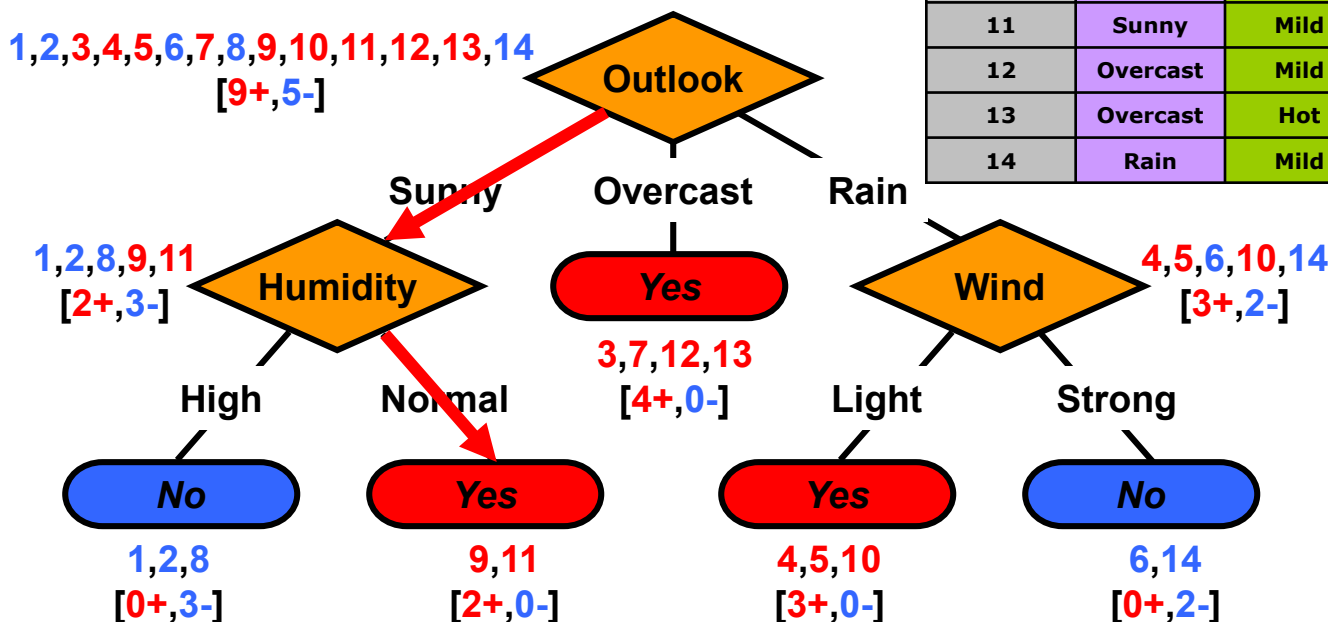
- 思想很简单，但经常很有效，至今在很多问题上，仍作为分类器使用
- 没有训练（建模）的过程：without training
- 属于非线性分类器
- 但，当标记样本数量很大、待处理对象维度很高时，计算复杂度很大；
- 也有很多值得研究的空间，如：如何降低计算复杂度、加权近邻法（距离远，权重小）尽管原理简单，

Decision Tree

■ Play Tennis

Day	Outlook	Temp.	Humidity	Wind	Play?
15	Sunny	Hot	Normal	Strong	Yes/No?

Day	Outlook	Temp.	Humidity	Wind	Play?
1	Sunny	Hot	High	Light	No
2	Sunny	Hot	High	Strong	No
3	Overcast	Hot	High	Light	Yes
4	Rain	Mild	High	Light	Yes
5	Rain	Cool	Normal	Light	Yes
6	Rain	Cool	Normal	Strong	No
7	Overcast	Cool	Normal	Strong	Yes
8	Sunny	Mild	High	Light	No
9	Sunny	Cool	Normal	Light	Yes
10	Rain	Mild	Normal	Light	Yes
11	Sunny	Mild	Normal	Strong	Yes
12	Overcast	Mild	High	Strong	Yes
13	Overcast	Hot	Normal	Light	Yes
14	Rain	Mild	High	Strong	No



Keys for constructing a DT:

- (1) splitting dimension
- (2) splitting value (for continuous dims)

ID3 [Quinlan, MLJ86]

CART [Brieman et al, Book84]

C4.5 [Quinlan, Book93]



Decision Tree

特点


- 得到的是一组规则集
- 决策过程具有良好的可理解性
- 对分类问题，在解决每类呈现多决策域分布、且交错分布的问题时，具有独特的优势
- 对于单一因素（特征）即可决定预测结果的情况，基于统计机器学习有时会出现的某些问题。如天气预测，如果湿度低到一定的值，就决定了肯定不会下雨；但在统计机器学习的思路下，如果其他多项条件都符合下雨条件，湿度因素就会被平均掉、忽略，从而会得出会下雨的错误预测。用决策树，则可以避免此类问题的出现。

dim 1



Support Vector Machine (SVM)


- How would you classify these points using a linear function in order to **minimize** the **error rate**?
- Linear function:
 - $f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$
 - A hyper-plane in the feature space
- **Infinite** number of answers!
- Which one is the **best**?



Linear SVM


Large Margin Linear Classifier

- The linear function with the **maximum margin** is the best
- **Margin** is defined as the width that the boundary could be increased before hitting a data point
- Why it is the best?
 - Robust to outliers and thus strong **generalization ability**



Non-Linear SVM

- Linearly separable:
- What if the data is too hard?
- How about...
 - mapping data to a high dimensional space




SVM




特点

- 在解决小样本、非线性问题上，具有独到的优势；因为对预测性能起关键因素的，是少数边界处的向量（支持向量）；只要边界处的向量分布正确、合理，预测效果就会较好
- 合适的变换核是关键，但对不同的问题，什么样的变换核有效，恰恰是个难点
- **LibSVM**：专门的平台

k -NN, Decision Tree and SVM



1-NN

Decision
TreeLinear
SVMNon-linear
SVM

Other Learners


- Naïve Bayesian（朴素贝叶斯）
- Neural Networks（神经网络）
- Least Squares（最小二乘）
- Gaussian Mixture Models
- Hidden Markov Models（隐马尔科夫模型，时序问题）
- Dynamic Bayesian Network（动态贝叶斯网络，时序问题）
- etc.

Different methods are suitable for different applications



Road Map

- Machine Learning
- Learners
- Learning Paradigms
- Resources
- Summary



Learning Paradigms (学习范式)

- **Ensemble learning** (集成学习)
- **Deep Learning** (深度学习)
- Semi-supervised learning (半监督学习)
- Cost-sensitive learning (代价敏感性学习)
- Class-imbalance learning (类别不平衡学习)
- Multi-label learning (多标记学习)
- Multi-instance learning (多示例学习)
-

Focus on ideas rather
than detailed algorithms



Ensemble Learning（集成学习）

An example: The gender recognition task

To predict whether a student is a boy or a girl



Useful features may include:

Height
Weight
Width of shoulder


Ensemble Learning（集成学习）

- Suppose 3 learners are trained with some collected data:
 - SVM, Neural Network(NN), Decision Tree(DT)
- For a test set of 100 students, the **best** accuracy of the three classifiers are all 90% (10 students are wrongly classified).
- If the classification results are as follows:
 - If only No.1-10 students are wrongly classified by SVM
 - If only No.11-20 students are wrongly classified by NN
 - If only No.21-30 students are wrongly classified by DT
- Fuse 3 classifiers by majority voting, we can get **100%** accuracy!



Ensemble Learning (集成学习)

- Training **multiple individual learners** for the same problem



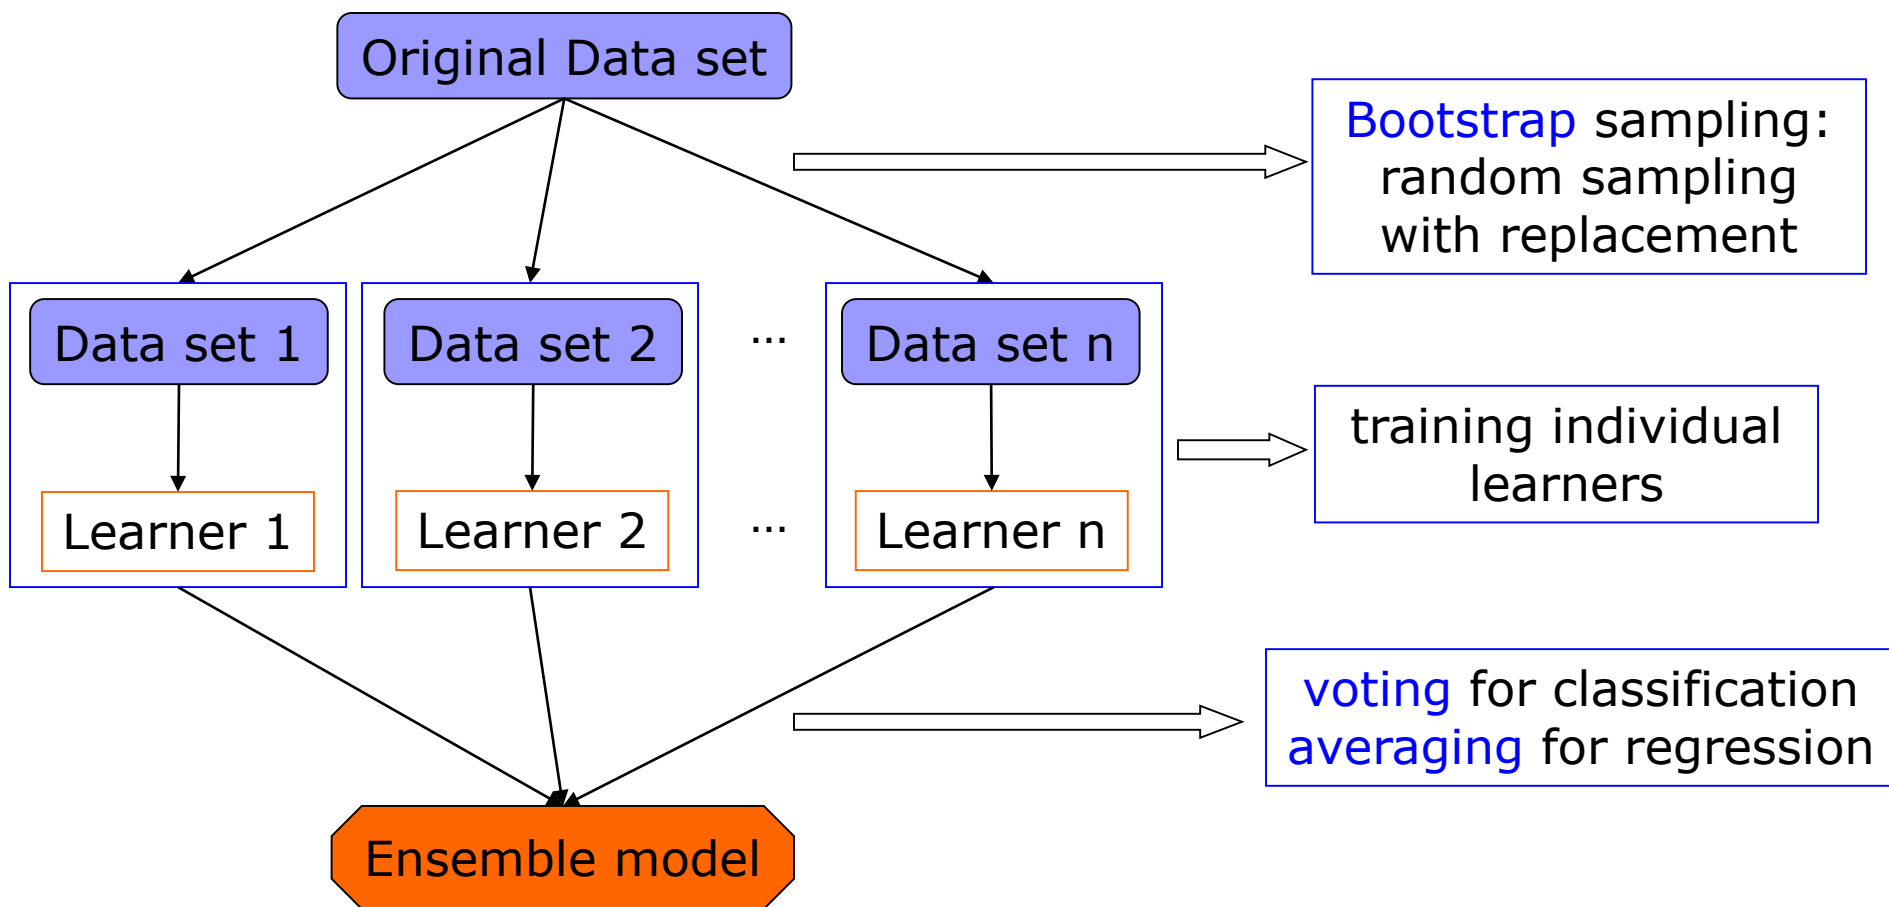
- Why ensemble?
 - The **generalization ability** of an ensemble is usually **significantly better** than the corresponding single learner

Base learners: the more accurate and the more diverse, the better

- Two steps:
 - (1) **train** base learners
 - (2) **combine** the individual predictions



Bagging(并行)

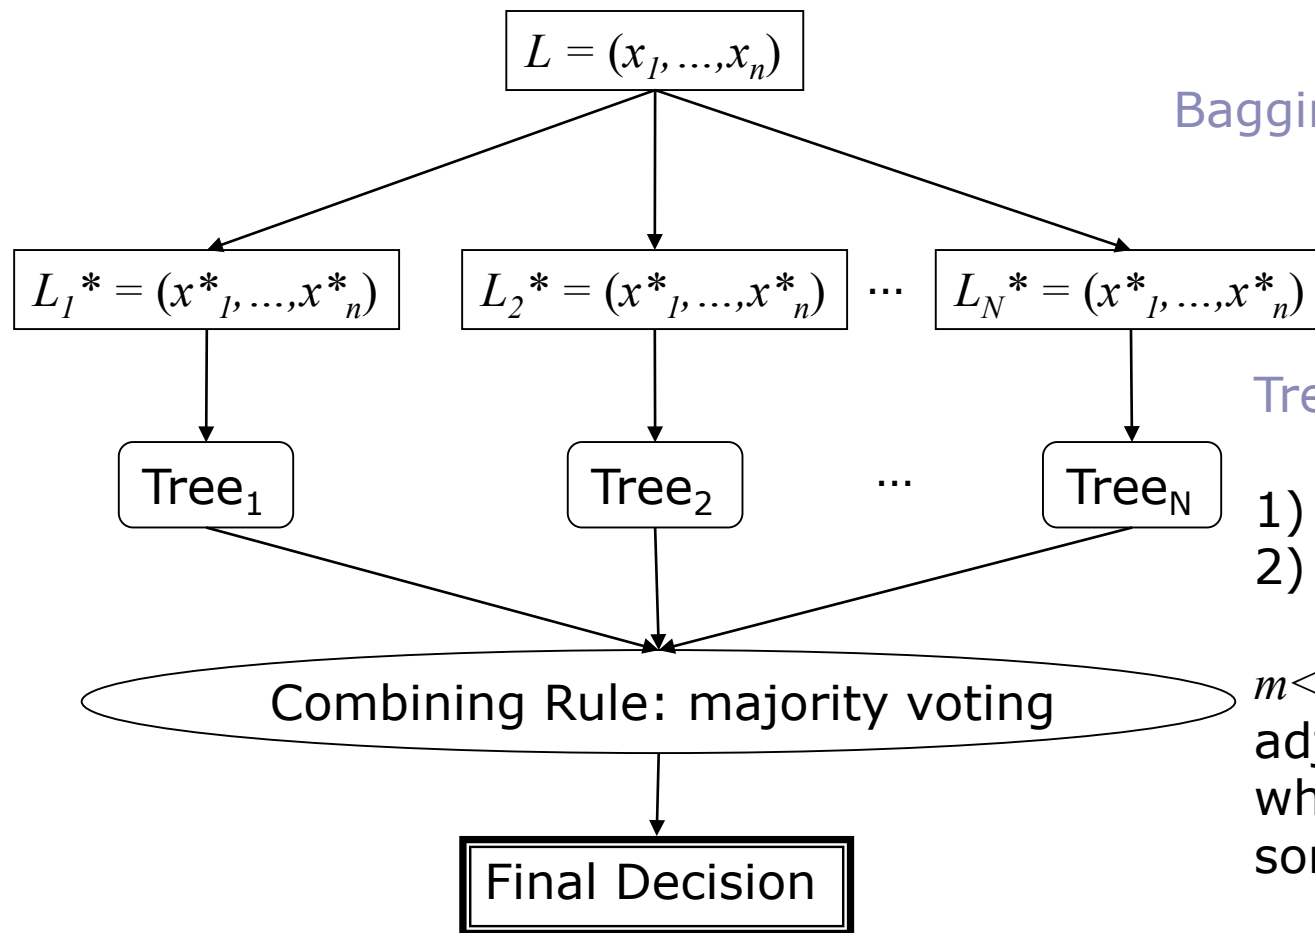


Learner → decision tree: Random Forest [Breiman, MLJ01]

Selective ensemble: Many Could be Better Than All [Zhou et al, AIJ02]



Bagging with decision trees: Random Forest



Bagging with decision trees

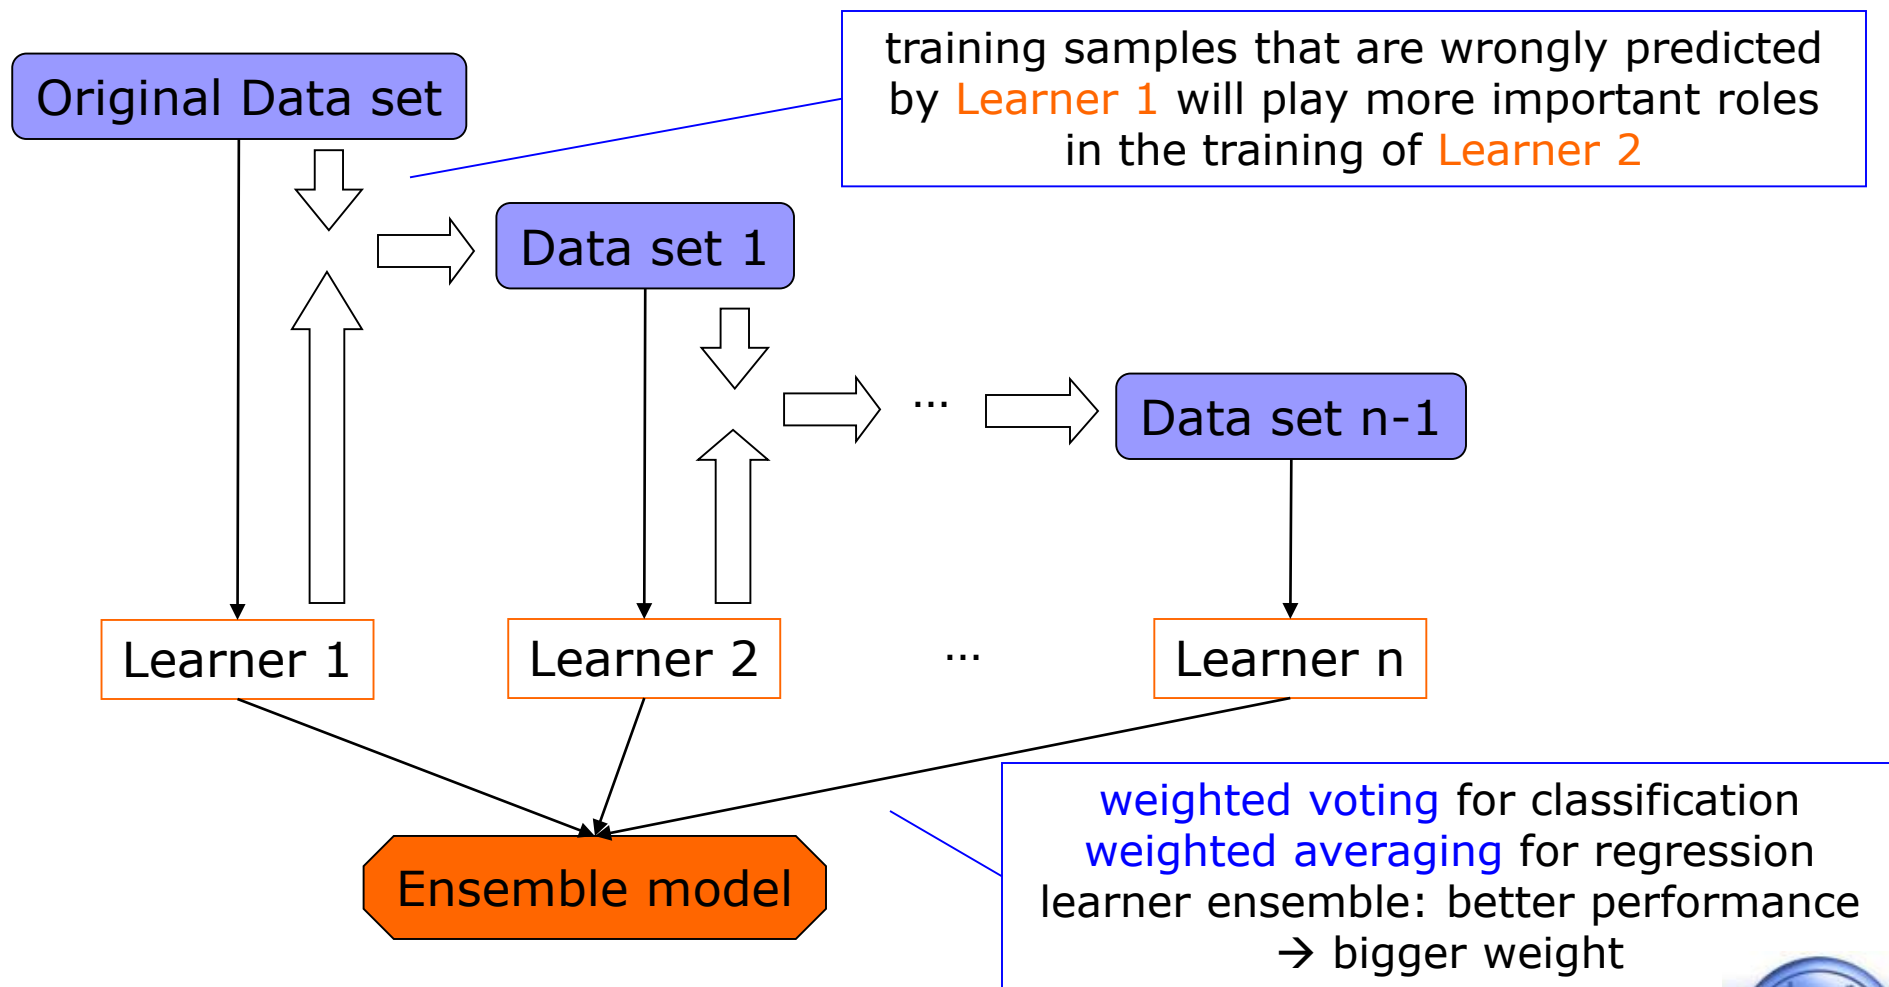
Tree growing process

- 1) no pruning;
- 2) m variables are used;

$m \ll M$: the **only** adjustable parameter to which random forests is somewhat sensitive



Boosting (串行)



Ensemble Learning（集成学习）

- Ensemble learning is not always helpful

Back to the gender recognition task(to 100 students):

If the three learners make the same mistakes（犯错同样的错误）

No.1-10 students are wrongly classified by SVM

No.1-10 students are wrongly classified by NN

No.1-10 students are wrongly classified by DT

Or if the performance of all three learners are poor（准确率都很低）

No.1-60 students are wrongly classified by SVM

No.11-70 students are wrongly classified by NN


No.21-80 students are wrongly classified by DT

- No improvement can be made by fusing the three learners.



Ensemble Learning (集成学习)

- The combined learners (classifiers) should have



How to get “good and different” individual learners is the key point of ensemble learning.



Ensemble Learning（集成学习）：


更为广义的理解

- 上面介绍的，是多学习器的集成，即，多个分类方法的集成
- 实际上，集成学习的思想，更为广义
- 还可以有更多不同水平的集成学习
 - Sensor level
 - Feature level
 - Score level
 - Decision level
 -
- 主要是看具体在哪个阶段进行集成



Ensemble Learning: at sensor level

- 通过不同的设备，能够获取人脸的可见光图像和红外图像（不同的采集设备、不同的传感器）
- 同时使用一个人的可见光图像和红外图像，实现人脸识别



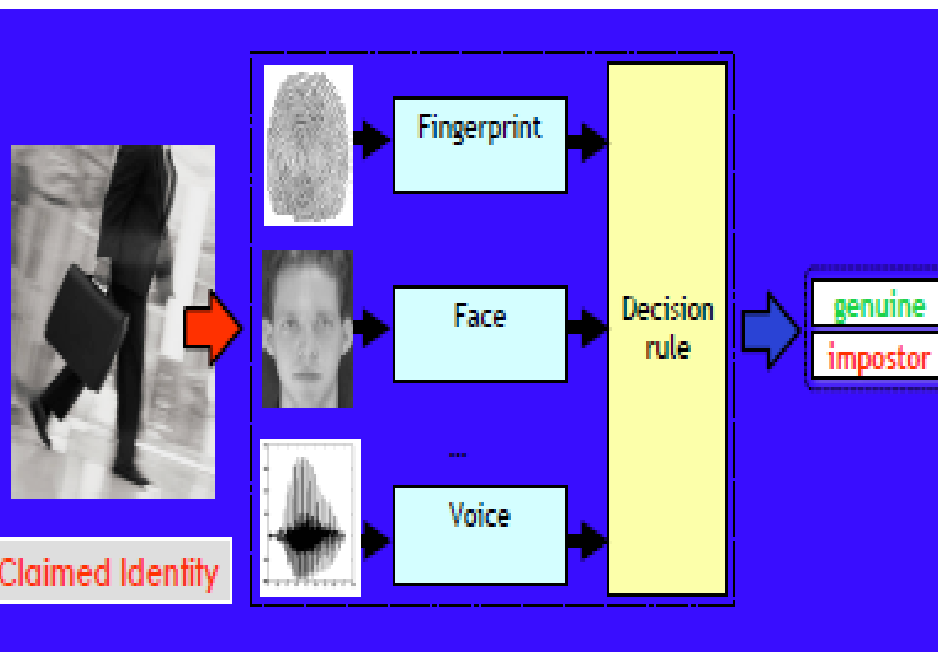
Ensemble Learning: at feature level(multi-view)

■ Multimodal biometrics

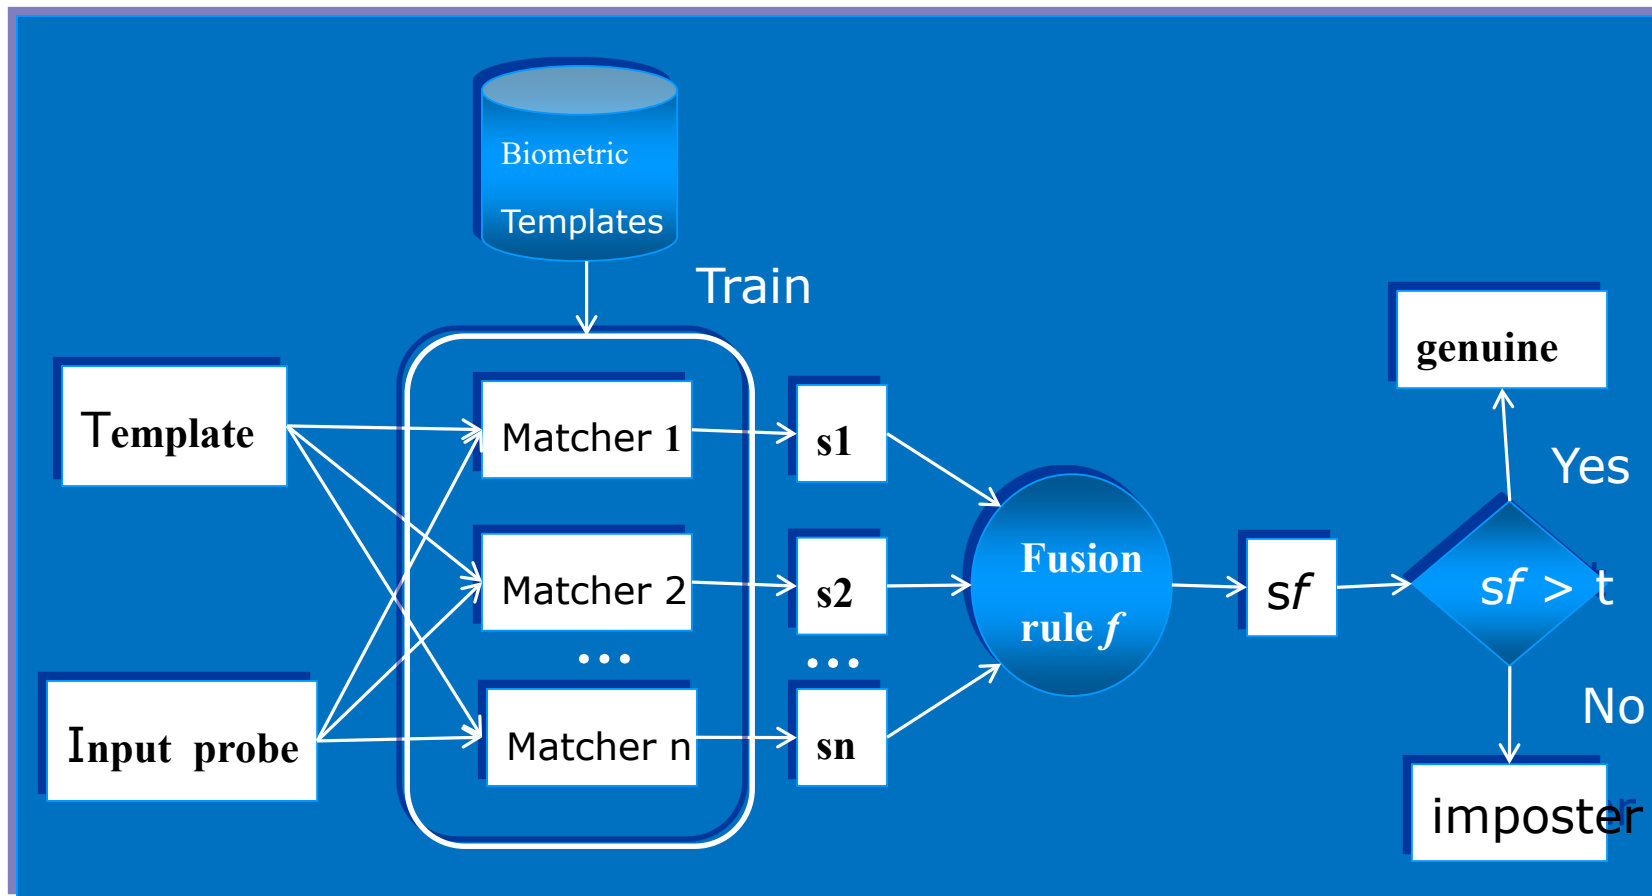
Fingerprint, face and voice
are three views of a person

■ Classification of Webpage

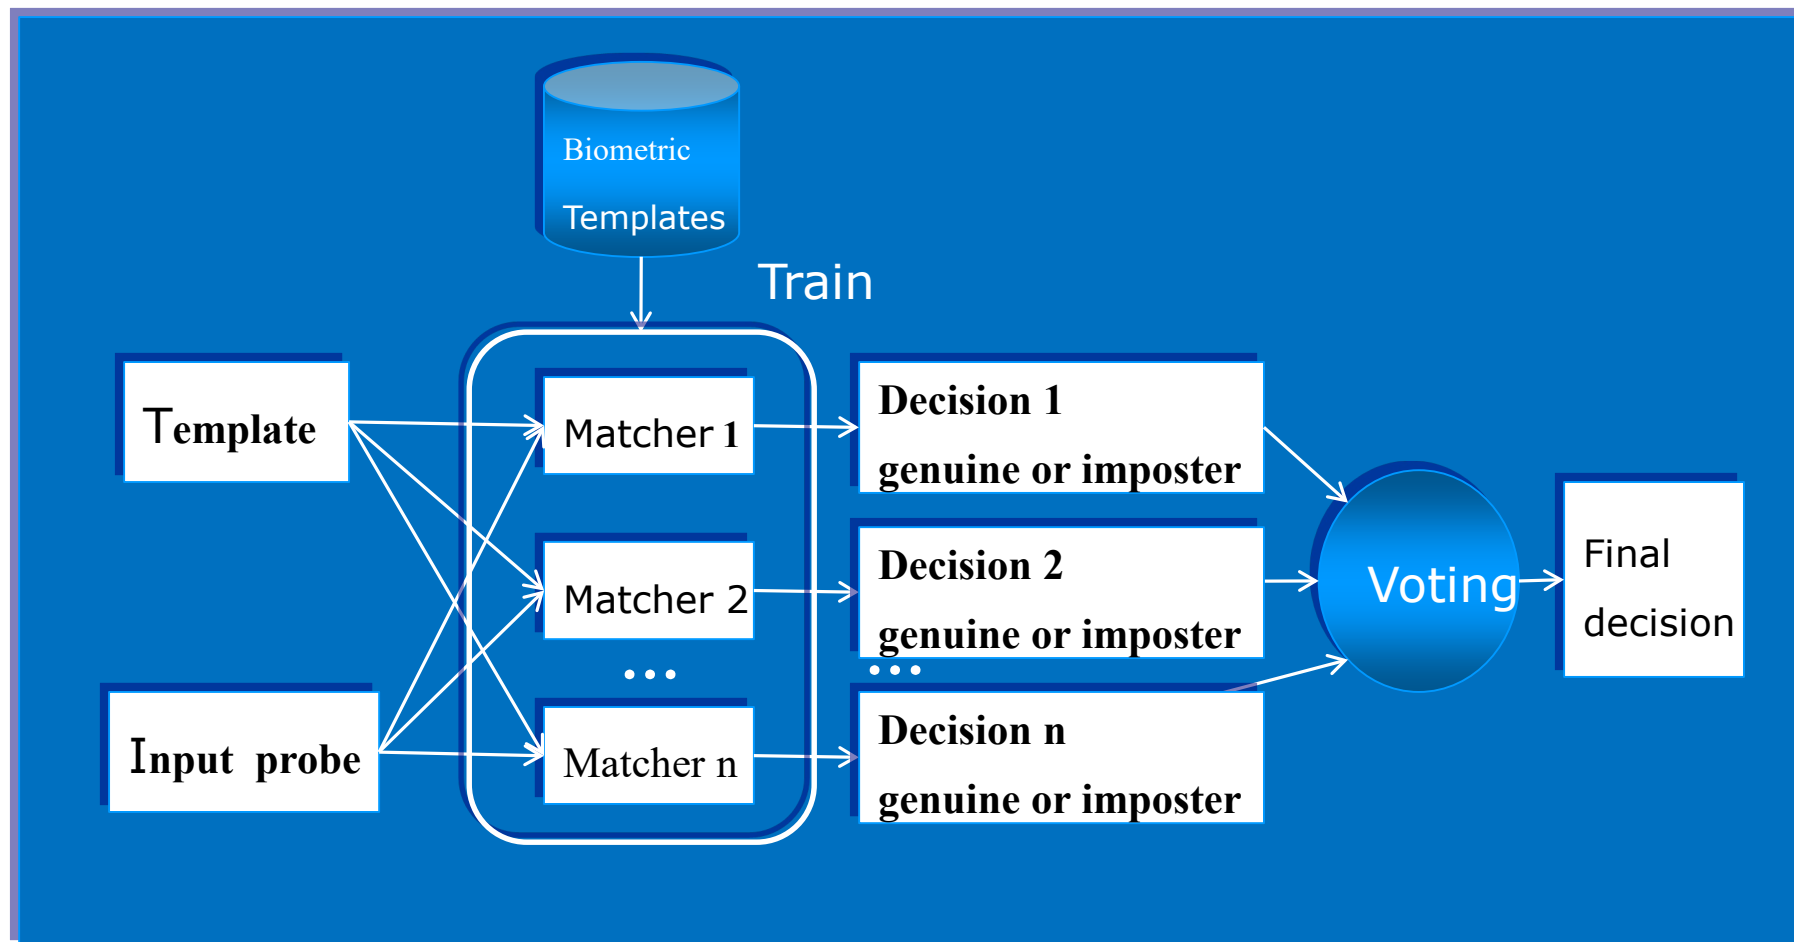
Texts, images and hyper-links
are three views of a webpage



Ensemble Learning: at score level



Ensemble Learning: at decision level



Ensemble Learning（集成学习）

- 思想很好理解（以多取胜、资源代价换取性能提升）
- （集成）程序很好实现
- 可用于各种预测问题：方法论、独立于具体问题
- 几乎适用于所有领域：万金油
- 从集成学习范式提出到深度学习出现（2006）之前，曾很长时间在机器学习领域独领风骚
- 效果如何，主要看两个方面：（1）被集成的基学习器性能如何；（2）被集成的多个学习器是否存在差异性（互补性）
- 对数学建模而言，可考虑预先准备多种学习器（决策树、支持向量机、神经网络等）的程序，根据建模任务的需要，现场集成（可尝试不同的集成策略）



集成学习方法的巨大成功

- ❑ KDDCup'07: 1st place for "... Decision Forests and ..."
- ❑ KDDCup'08: 1st place of Challenge1 for a method using Bagging; 1st place of Challenge2 for "... Using an Ensemble Method "
- ❑ KDDCup'09: 1st place of Fast Track for "Ensemble ... "; 2nd place of Fast Track for "... bagging ... boosting tree models ...", 1st place of Slow Track for "Boosting ... "; 2nd place of Slow Track for "Stochastic Gradient Boosting"
- ❑ KDDCup'10: 1st place for "... Classifier ensembling"; 2nd place for "... Gradient Boosting machines ... "



集成学习方法的巨大成功

- ❑ KDDCup'11: 1st place of Track 1 for "A Linear Ensemble ..."; 2nd place of Track 1 for "Collaborative filtering Ensemble"; 1st place of Track 2 for "Ensemble ..."; 2nd place of Track 2 for "Linear combination of ..."
- ❑ KDDCup'12: 1st place of Track 1 for "Combining... Additive Forest..."; 1st place of Track 2 for "A Two-stage Ensemble of..."
- ❑ KDDCup'13: 1st place of Track 1 for "Weighted Average Ensemble"; 2nd place of Track 1 for "Gradient Boosting Machine"; 1st place of Track 2 for "Ensemble the Predictions"



集成学习方法的巨大成功

- ❑ KDDCup'14: 1st place for “ensemble of GBM, ExtraTrees, Random Forest...” and “the weighted average”; 2nd place for “use both R and Python GBMs”; 3rd place for “gradient boosting machines... random forests” and “the weighted average of...”
- ❑ KDDCup'15: 1st place for “Three-Stage Ensemble and Feature Engineering for MOOC Dropout Prediction”
- ❑ Netflix Prize:
 - ✓ 2007 Progress Prize Winner: Ensemble
 - ✓ 2008 Progress Prize Winner: Ensemble
 - ✓ 2009 \$1 Million Grand Prize Winner: Ensemble !!



Ensemble Learning: Tutorials

- L Rokach, Pattern Classification Using Ensemble Methods, World Scientific, Singapore, 2010
- Z-H Zhou, Ensemble Methods: Foundations and Algorithms, Chapman & Hall/CRC, Boca Raton, 2012
- Z-H Zhou, Machine Learning, Chapter 8, 2016



Ensemble Learning:

Useful web pages and tools

■ Web pages:

- http://www.scholarpedia.org/article/Ensemble_learning
- https://en.wikipedia.org/wiki/Ensemble_learning
- <https://beta.learning.intersystems.com/course/view.php?id=15>

■ Online Tools

- The Waffles (machine learning) toolkit
[https://en.wikipedia.org/wiki/Waffles_\(machine_learning\)](https://en.wikipedia.org/wiki/Waffles_(machine_learning))



Deep Learning（深度学习）

当前最火的机器学习范式，在图像分类、物体检测与识别、语音识别等领域，取得了突破性进展

- 人工神经网络
- 多隐层人工神经网络
- 什么是深度学习
- 深度学习的发展历史
- 取得的突破性进展
- 为何有效
- 什么情况下应该考虑使用
- 发展趋势和局限
- 有用资源




Artificial Neural Network (ANN)

Our brains are a huge network of processing elements. A typical brain contains a network of 10 billion neurons.




Biological Neuron and ANN

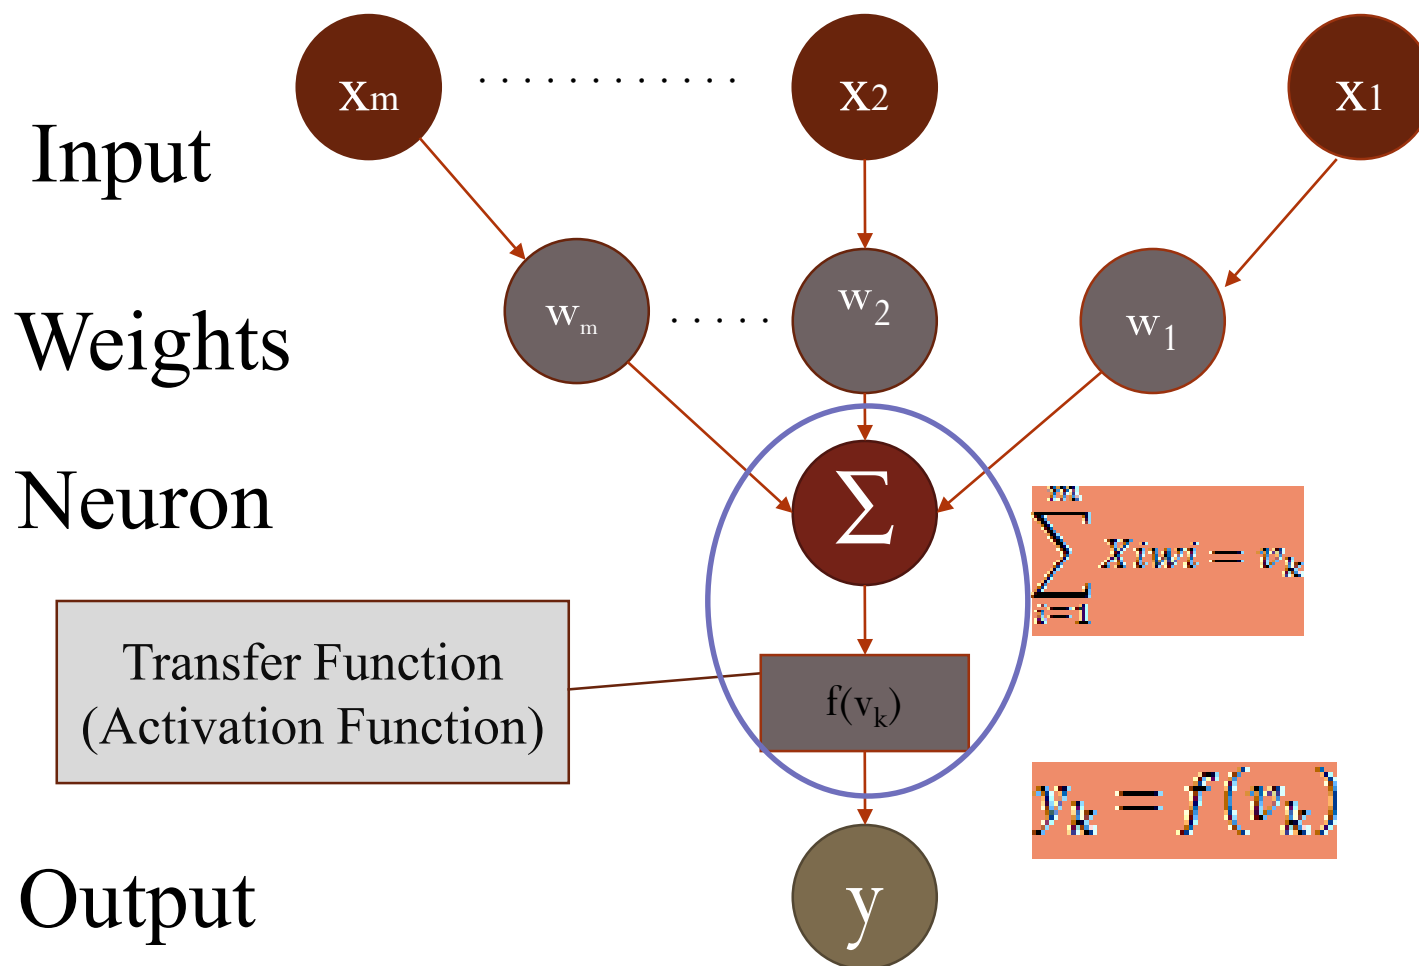
Biological Neuron




Artificial Neuron




ANN--单层感知器



ANN—单隐层神经网络



ANN—多隐层神经网络



Deep Learning（深度学习）：

What is it

- 本质上，就是多隐层人工神经网络
- 属于深层模型（**Deep Model**），SVM等常见学习器都属于浅层模型（**Shallow Model**）
- 针对具体任务，利用给定的一批标记数据，先训练一个多隐层神经网络，然后使用它，这就是深度学习
- 深度学习与浅层模型几个突出的不同点
 1. 自动学习特征 Vs. 经验知识+人工定义特征
 2. 端到端（**end to end**）Vs. 分步、分治
 3. 超强的非线性建模能力 Vs. 有限的非线性建模能力



Deep Learning（深度学习）：

History

- 与多隐层神经网络有关的几个重要时间节点
 - 上世纪80年代就出现了多隐层人工神经网络
 - Before 2000（受到的关注少）
 - 2000-2005（Bengio等，在推动深度学习，但在应用上取得的进展有限）
 - 2006至今（Hinton等在《科学》发表论文：优异特征学习能力、无监督的分层预训练）：进入深度学习的时代
 - 优化方法、优化技巧在近几年深度学习的发展中起了很重要的作用




Deep Learning（深度学习）：

Why it could not work well before 2006

- 设备条件限制：缺乏高性能计算装备，更没有GPU
- 数据条件限制：缺乏大数据（训练数据）
- 技术条件限制：对多隐层神经网络，缺乏有效的训练方法
- 2006年以后，上述限制逐步解除了



突破性进展1: 图像分类: on ImageNet




- Database:
1000 categories,
1.2 million training
images,
150,000 testing
images.
- Task:
classify testing image
into one of
1000 categories.



突破性进展1:

图像分类: on ImageNet



突破性进展2:

人脸识别: on LFW

- LFW是用于人脸识别的一个相对高难度、大尺寸公开数据库
- 基于深度学习的DeepID人脸识别技术在LFW库上的准确率达到99.77%，比人眼识别更加精准！



突破性进展3:

深度学习在中文语音识别上超过了人类

- 2015年12月，百度研究院硅谷人工智能实验室
- 中文语音测试：人类语音识别的错误率是4.0%，而机器是3.7%
- 百度首席科学家吴恩达：对于无上下文的短语，基于深度学习的计算机系统的识别能力超过了人类



突破性进展4:

围棋人机大战: AlphaGo Vs. 李世石

- 在上世纪90年代末期以前, 普遍认为: 尽管机器可在国际象棋比赛中战胜人类棋手, 但机器永远不可能, 至少在可以预见的很长时间内, 机器不可能在围棋上战胜人类
- 2016年3月9-15日, AlphaGo以4:1战胜韩国九段棋手李世石
- 当前世界围棋界, 应该没有人可以战胜AlphaGo了, 今后, 机器的优势将更加突出
- 深度学习, 是AlphaGo使用的核心技术之一



Deep Learning（深度学习）：

最新进展

- Attention（注意力机制）
- Reasoning（推理）
- Planning and Reinforcement Learning（规划与强化学习）
- 当前，最前沿的是融合技术。比如，视觉与自然语言理解的结合
- 向后看，可能的前沿方向：无人驾驶、推理和回答问题



Deep Learning（深度学习）：

为何展示出极其突出的性能优势？

- 具有自动学习特征的能力（Feature Learning）
- 学习到的特征体系和人工定义特征不同，在完备性和非冗余性上，更准确地说，在区分性上，强于人工定义的特征
- 对复杂分类问题，有能力学习到极其复杂的“分界面”（过拟合，非贬义）
- 解决问题的思路和技术框架，都具有较强的通用性



Deep Learning（深度学习）：

什么情况下应该考虑使用深度学习

- 针对任务要求，难以人工定义特征（因为这需要先验知识），或者人工定义的特征不够有效：**Feature Learning**
- 大量标记样本（期望训练集分布更接近全集分布）。但对标记样本数量的要求也不是那么绝对（很多时候，使用少量标记样本也很有效，核心还是分布问题）
- 高性能计算资源：**GPU**（尤其是对图像、语音）；但这是针对训练过程；应用时，普通电脑就够了
- 有效的网络训练方法（这是对科研而言，对竞赛，使用已有的、相对成熟的网络训练方法即可）
- 面向数学建模：（1）使用开源框架和相近数据，预先训练网络；（2）竞赛现场，利用给定数据（数据量一般不大），进行再训练（训练量应该不大），这是迁移学习的概念



Deep Learning:

有用资源

► <http://blog.csdn.net/zouxy09/article/details/8775360>

了解一些deep learning基本方法的思想

► <http://ufldl.stanford.edu/wiki/index.php/UFLDL>
教程

deep learning大牛Andrew Ng所写，还有实验、源代码，推荐细读



Deep Learning:

有用资源

■ Platforms:

- Pytorch: <https://pytorch.org/>
- TensorFlow: <https://tensorflow.google.cn/>
- Keras: <https://keras.io/>





Semi-Supervised Learning

- Why semi-supervised?
 - $\mathbf{X}_{1:n}$ known, $\mathbf{Y}_{1:l}$ known, $\mathbf{Y}_{l+1:n}$ unknown
 - **labeled** data may be hard to get; **unlabeled** data is cheap
 - people want better performance for **free**
- Example: Web Page Classification




Learnability


- **Goal:** using both labeled and unlabeled data to build better learners, rather than using each one alone.




labeled data



GMM model



adding
unlabeled data



GMM model


Semi-Supervised Learning Algorithms

- Self-training
- Generative models
 - Gaussian mixture models + EM
- Semi-Supervised SVM (**S3VM**)
 - a.k.a.: Transductive SVM (TSVM)
- Disagreement-based algorithms
 - **Co-training**
 - Tri-training
- Graph-based algorithms
 - Label propagation
 - **Manifold regularization**
 - local and global consistency



Cost-Sensitive Learning

- Why cost-sensitive learning?
 - Traditional view: low **error rate** → good performance
 - However, in many real applications, different mistakes often have different **costs**
 - We should minimize the **total cost** instead of simply minimizing the **error rate**



Keys: (1) estimate the misclassification costs;
(2) minimize the total cost;

CSL Algorithms


- Direct modification on traditional learners
 - cost-sensitive SVM
 - cost-sensitive decision tree
 - cost-sensitive neural networks
 - cost-sensitive boosting

- Rescaling
 - re-weighting
 - re-sampling
 - undersampling & oversampling
 - threshold moving
 - MetaCost



Class-Imbalance Learning


- Why class-imbalance learning?
 - In many real applications, the data sets are typically **imbalanced**, i.e., some classes have much more instances than others.
- Cancer detection
 - healthy : cancer = 99:1
 - minimizing **error rate**:
classify all instances as healthy → **1%** error rate
- CIL algorithms
 - **cost-sensitive learning**: re-sampling, re-weighting, etc.
 - **one-class learning**: one-class SVM




Multi-Label Learning

■ Why multi-label?

- In **traditional** supervised learning: a real-world **object** is represented by **an instance**
- The instance is associated with **a label** which indicates the concerned characteristics of the object



but



Elephant?

Lion?

Bush(丛林)?

Tropic(热带)?

Africa?

Multi-label learning: **an object** is attached with **multiple labels**

MLL Algorithms

- **Decomposing** the task into multiple binary classification problems each for a class
 - MLSVM
- Considering the **ranking** among labels
 - BoosTexter
 - BP-MLL
 - RankSVM
- Exploring the **label correlation**
 - Probabilistic generative models
 - Maximum entropy methods

Elephant?
Lion?
Bush?

rank1

Tropic?
Africa?

rank2



Penguin(企鹅)?

Iceberg(冰山)?



Multi-Instance Learning

- Why multi-instance?
 - Multi-label learning only addresses the **output ambiguity**
 - How about the **input ambiguity**?



total image \Rightarrow a bag

regions in image \Rightarrow instances in bag

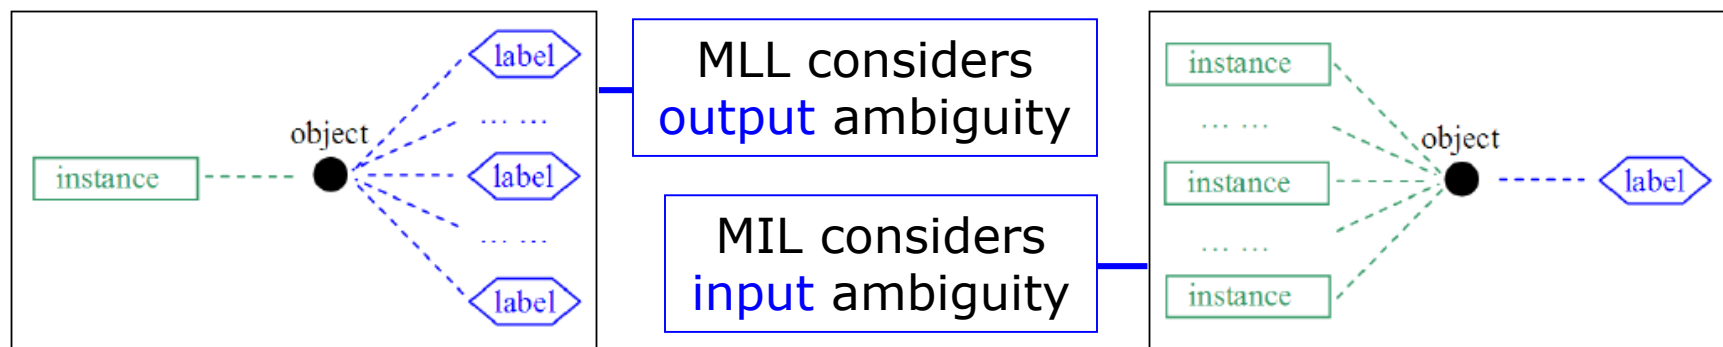
Multi-instance learning:
 an **object** contains **multiple instances**

MIL Algorithms

- Direct modification on traditional learners
 - by shifting their focuses from the discrimination on the **instances** to the discrimination on the **bags**
 - k -NN \rightarrow Bayesian- k NN, Citation- k NN
 - decision tree \rightarrow Relic, ID3-MI, RIPPER-MI
 - SVM \rightarrow MI-SVM, mi-SVM, DD-SVM
- Other topics
 - density estimation: Diverse Density, EM-DD
 - kernel computation: multi-instance kernels
 - regression: MI-LR
 - clustering: BAMIC
 - ensemble: MI-Ensemble, MI-Boosting



Multi-Instance Multi-Label Learning



Input and output ambiguities usually occur simultaneously!

tree

$[a_1, a_2, \dots, a_d]$

elephant

$[b_1, b_2, \dots, b_d]$


lion

$[c_1, c_2, \dots, c_d]$

.....

grass

$[u_1, u_2, \dots, u_d]$




Elephant?

Lion?

Bush?

Tropic?

Africa?



Multi-instance multi-label learning (MIML):

each object contains **many instances** and is attached with **multiple labels**

Other Learning Paradigms

- Learning to Rank
[Liu, FTIR09] & **Hang Li**
- Online Learning & Incremental Learning
[Shwartz, Thesis07] & **Yoram Singer**
- Transfer Learning
[Pan & Yang, TKDE (in press)] & **Qiang Yang**




Other Learning Paradigms

- Multi-Task Learning
[Evgeniou et al, JMLR05] & [Argyriou et al, NIPS'06] & Andreas Argyriou
- Reinforcement Learning
[Kaelbling et al, JAIR96] & [Sutton & Barto, Book98]
- Active Learning
[Tong, Thesis01]
- etc.



Road Map

- Machine Learning
- Learners
- Learning Paradigms
- Resources
- Summary




Resources: Books

- 周志华著，《机器学习》，中文，清华大学出版社，ISBN号：978-7-302-42328-7，定价：88元，2016年1月出版（特别推荐）



Resources: Books

- Ian Goodfellow, Yoshua Bengio, Aaron Courville. [Deep Learning](#), 2017.



Resources: Books

- R. Duda, P. Hart, D. Stork: [Pattern Classification](#), 2nd Edition, Wiley, 2000 (入门书)
- C. Bishop: [Pattern Recognition and Machine Learning](#), Springer, 2006 (难度大一些)
- T. Mitchell: [Machine Learning](#), McGraw Hill, 1997 (老一些)



Top Conferences/Top Journals

- Top Conferences

- ICML, COLT, NIPS, ACML, etc. (侧重机器学习理论)
- IJCAI, AAAI, ICCV, CVPR, ECCV, etc. (侧重机器学习应用)

- Top Journals

- JMLR, AI, TPAMI, IJCV, TKDD, TKDE, PR, etc.



Resources: International Scholars

- **Tom Mitchell**: <http://www.cs.cmu.edu/~tom/>
- **Michael Jordan**: <http://www.cs.berkeley.edu/~jordan/>
- **Geoffrey Hinton**: <http://www.cs.toronto.edu/~hinton/>
- **Bernhard Schölkopf**: <http://www.kyb.mpg.de/~bs/>
- **Alexander Smola**: <http://alex.smola.org/>

- **Rong Jin**: <http://www.cse.msu.edu/~rongjin/>
- **Jieping Ye**: <http://www.public.asu.edu/~jye02/>
- **Tong Zhang**: <http://www.stat.rutgers.edu/~tzhang/>
- **Andrew Ng**: <http://ai.stanford.edu/~ang/>
- **Eric Xing**: <http://www.cs.cmu.edu/~epxing/>
- **Fei Sha**: <http://www-rcf.usc.edu/~feisha/>
- **Xiaojin Zhu**: <http://pages.cs.wisc.edu/~jerryzhu/>



Resources: Tools


■ Developing tools

- **MATLAB**: <http://www.mathworks.com/>
- **WEKA**: <http://www.cs.waikato.ac.nz/~ml/weka/>
- **LIBSVM**: <http://www.csie.ntu.edu.tw/~cjlin/libsvm/>
- **MOSEK**: <http://www.mosek.com/>

■ Data sets

- **UCI ML Repository**: <http://archive.ics.uci.edu/ml/>
- **UCI KDD Archive**: <http://kdd.ics.uci.edu/>
- **Clustering data**: <http://cs.joensuu.fi/sipu/datasets/>

■ Google and WIKIPEDIA




!



Road Map

- Machine Learning
- Learners
- Learning Paradigms
- Resources
- Summary



Summary

- Machine learning
 - concept
 - applications
- Learners
 - k -NN
 - decision tree
 - SVM
- Learning paradigms
 - ensemble learning
 - deep learning



References

- [Mitchell, Book97] T. Mitchell. Machine Learning. McGraw Hill, New York, 1997.
- [Quinlan, MLJ86] J. Quinlan. Induction of Decision Trees. Machine Learning, 1(1): 81-106, 1986.
- [Breiman et al, Book84] L. Breiman, J. Friedman, R. Olshen and C. Stone. Classification and Regression Trees. Chapman & Hall, New York, 1984.
- [Quinlan, Book93] J Quinlan. C4.5: Programs for Machine Learning. Morgan Kaufmann, San Francisco, CA, 1993.
- [Vapnik, Book95] V. Vapnik. The Nature of Statistical Learning Theory. Springer, New York, 1995.
- [Zhou, EB09] Z.-H. Zhou. Ensemble Learning. In: Encyclopedia of Biometrics. Springer, Berlin, 270-273, 2009.
- [Breiman, MLJ96] L. Breiman. Bagging Predictors. Machine Learning, 24(2): 123-140, 1996.
- [Breiman, MLJ01] L. Breiman. Random Forests. Machine Learning, 45(1): 5-32, 2001.



References

- [Zhou et al, AIJ02] Z.-H. Zhou, J. Wu and W. Tang. Ensembling Neural Networks: Many Could Be Better Than All. Artificial Intelligence 137(1-2): 239–263, 2002.
- [Freund & Schapire, JCSS97] Y. Freund and R. Schapire. A Decision-Theoretic Generalization of Online Learning and an Application to Boosting. Journal of Computer and System Sciences, 55(1): 119–139, 1997.
- [Zhu, TR08] X. Zhu. Semi-Supervised Learning Literature Survey. Computer Sciences TR 1530, University of Wisconsin Madison, 2008.
- [Ling & Sheng, EML08] C. Ling and V. Sheng. Cost-Sensitive Learning and the Class Imbalance Problem. In: Encyclopedia of Machine Learning. Springer, Berlin, 2008.
- [Weiss, KDDEXP04] G. Weiss. Mining With Rarity: A Unifying Framework. SIGKDD Explorations, 6(1): 7-19, 2004.
- [Tsoumakas & Katakis, IJDWM07] G. Tsoumakas and I. Katakis. Multi-Label Classification: An Overview. International Journal of Data Warehousing and Mining, 3(3): 1-13, 2007.



References

- [Tsoumakas et al, DMKDH10] G. Tsoumakas, I. Katakis and I. Vlahavas. Mining Multi-label Data. In: Data Mining and Knowledge Discovery Handbook, 2nd edition, Springer, Berlin, 2010.
- [Zhou, JSCT06] Z.-H. Zhou. Multi-Instance Learning from Supervised View. Journal of Computer Science and Technology, 21(5): 800-809, 2006.
- [Zhou & Zhang, NIPS'06] Z.-H. Zhou and M.-L. Zhang. Multi-Instance Multi-Label Learning with Application to Scene Classification. Advances in NIPS, 1609-1616, 2006.
- [Liu, TR06] Y. Liu. Distance Metric Learning: A Comprehensive Survey. Technical Report, Michigan State University, 2006.
- [Shwartz, Thesis07] S. Shalev-Shwartz. Online learning: Theory, Algorithms, and Applications. PhD Thesis, Hebrew university, 2007.
- [Bengio, FTML09] Yoshua Bengio. Learning Deep Architectures for AI. Foundations and Trends in Machine Learning, 2(1): 1-127, 2009.
- [Pan & Yang, TKDE2010] S. Jialin Pan and Q. Yang. A Survey on Transfer Learning. IEEE Transactions on Knowledge and Data Engineering.



References

- [Liu, FTIR09] T.-Y. Liu. Learning to Rank for Information Retrieval. Foundations and Trends in Information Retrieval, 3(3): 225-331, 2009.
- [Evgeniou et al, JMLR05] T. Evgeniou, C. Micchelli and M Pontil. Learning Multiple Tasks with Kernel Methods. Journal of Machine Learning Research, 6:615-637, 2005.
- [Argyriou et al, NIPS'06] A. Argyriou, T. Evgeniou and M. Pontil. Multi-Task Feature Learning. Advances in NIPS, 41-48, 2007.
- [Kaelbling et al, JAIR96] L. Kaelbling, L. Michael and M. Andrew. Reinforcement Learning: A Survey. Journal of Artificial Intelligence Research, 4:237-285, 1996.
- [Sutton & Barto, Book98] R. Sutton and A. Barto. Reinforcement Learning: An Introduction. MIT Press, Cambridge, MA, 1998.
- [Tong, Thesis01] S. Tong. Active Learning: Theory and Applications. PhD Thesis, Stanford University, 2001.

